

# Advanced RNNs

NLP  
Andreas Marfurt

# Overview

- Bidirectional RNNs
- Using RNNs for
  - Classification
  - Generation
  - Sequence-to-sequence tasks
- Attention in RNNs
- Contextual Word Embeddings: ELMo

# Advanced RNN Topics

# Bidirectionality

- In RNNs, we go through our input sequence from left to right
  - For right-to-left (RTL) languages, the input sequence is simply reversed
- Consider the example sentence:

“The bank consists of the sides of the river.”

# Bidirectionality

- In RNNs, we go through our input sequence from left to right
  - For right-to-left (RTL) languages, the input sequence is simply reversed
- Consider the example sentence:

What type of bank?

—————→  
“The bank consists of the sides of the river.”

# Bidirectionality

- In RNNs, we go through our input sequence from left to right
  - For right-to-left (RTL) languages, the input sequence is simply reversed
- Consider the example sentence:

What type of bank?

“The bank consists of the sides of the river.”


It must be a river bank ← The topic is rivers

# Bidirectionality

“The bank consists of the sides of the river.”


- The right context of bank has information that the left context doesn't have.
- Idea: Use both context directions.

# Bidirectional RNN (BiRNN)







# Bidirectional RNN (BiRNN)




# Bidirectional RNN (BiRNN)



# Bidirectional RNN (BiRNN)



# Bidirectional RNN (BiRNN)



Multiple layers:  
concatenate  
 $[\overrightarrow{h}_1^1; \overleftarrow{h}_1^1]$ ,  
give as input to  
next layer

# BiRNN Definition

- RNN:

$$h_t = \sigma(W_x x_t + W_h h_{t-1} + b_h)$$
$$y_t = \text{softmax}(W_y h_t + b_y)$$

- BiRNN:

$$\vec{h}_t = \sigma(\vec{W}_x x_t + \vec{W}_h \vec{h}_{t-1} + \vec{b}_h)$$
$$\overleftarrow{h}_t = \sigma(\overleftarrow{W}_x x_t + \overleftarrow{W}_h \overleftarrow{h}_{t-1} + \overleftarrow{b}_h)$$
$$y_t = \text{softmax}(W_y [\vec{h}_t; \overleftarrow{h}_t] + b_y)$$

# BiRNN Definition

- RNN:

$$h_t = \sigma(W_x x_t + W_h h_{t-1} + b_h)$$
$$y_t = \text{softmax}(W_y h_t + b_y)$$

- BiRNN:


$$\vec{h}_t = \sigma(\vec{W}_x x_t + \vec{W}_h \vec{h}_{t-1} + \vec{b}_h)$$
$$\overleftarrow{h}_t = \sigma(\overleftarrow{W}_x x_t + \overleftarrow{W}_h \overleftarrow{h}_{t-1} + \overleftarrow{b}_h)$$
$$y_t = \text{softmax}(W_y [\vec{h}_t; \overleftarrow{h}_t] + b_y)$$

Multiple layers:  
Input becomes  
 $[\vec{h}_1^{(l-1)}; \overleftarrow{h}_1^{(l-1)}]$   
(adjust input dim.  
to 2x hidden dim.)

# Classification


- Final hidden state  $h_n$   
“encodes” all information  
of the input sequence

→ Use  $h_n$  as a feature for  
classification



# Sequence Classification


- Classify from  $h_n$ :
  - Acceptability: Is this sentence grammatical?
- Compare two sentences (e.g. to determine if one is a paraphrase of the other):
  - Get final hidden states of both sentences:  $h_n, h'_m$
  - Compute similarity (e.g. cosine similarity)





# Token Classification

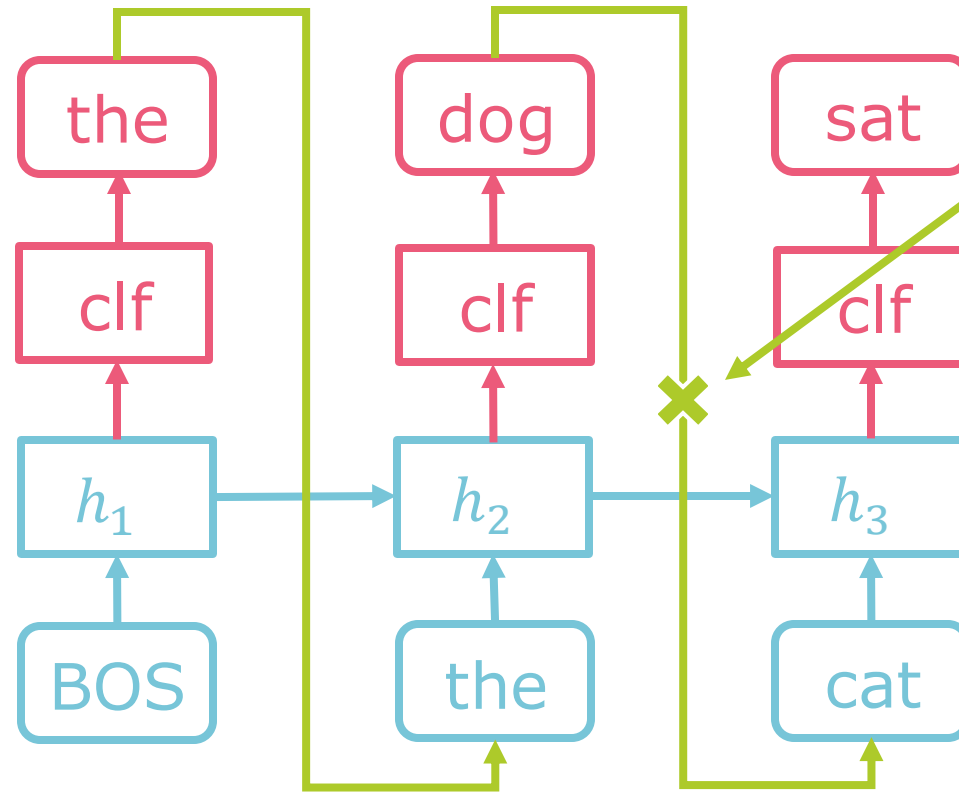
- One decision for each input token
  - PoS tagging
  - Named entity recognition
- ... or text generation!



# Text Generation

- Text generation is also a token classification task:
  - Next word prediction is just a probability distribution over the vocabulary
  - One strategy: Choose the vocabulary item with the highest probability
    - This is called *greedy decoding*
    - We will see later why this isn't always a good idea

# Text Generation



*Teacher forcing:*  
During training, the true word is given to the network as the next input, irrespective of what it predicted.

# Sequence-to-sequence (seq2seq)

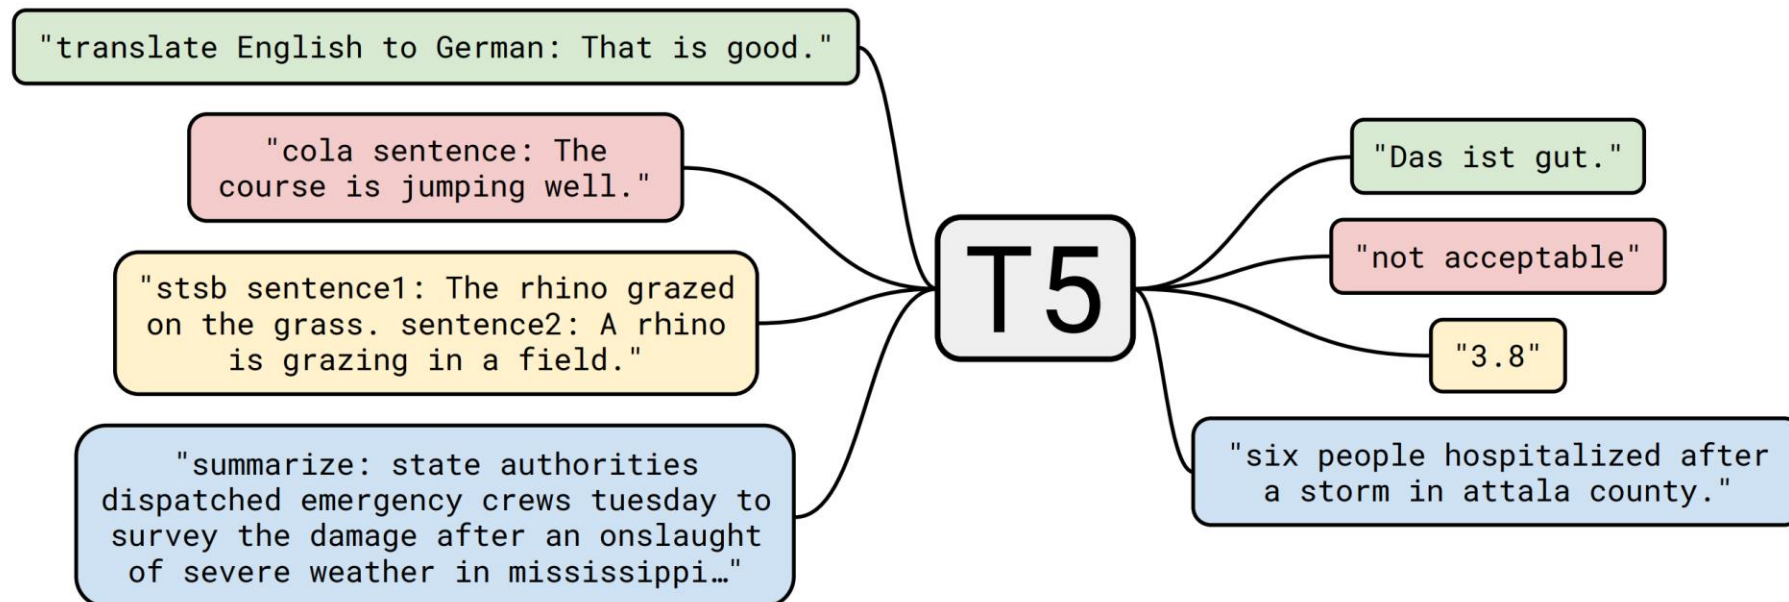
- Sequence-to-sequence tasks are very popular right now
  - Sequence as input
  - Sequence as output
  - Network does the transformation
- Tasks:
  - Machine translation
  - Summarization
  - Question answering
  - ... and nearly every task: Just formulate it in a text-to-text format.

# Sequence-to-sequence (seq2seq)

- Arithmetic operations
  - Input:  $1 + 2 = ?$
  - Output: "3" (as a string)


# Sequence-to-sequence (seq2seq)

- Arithmetic operations
  - Input:  $1 + 2 = ?$
  - Output: "3" (as a string)




# Sequence-to-sequence (seq2seq)

- Encoder:
  - Understands the text
  - Extracts the important information
  - Outputs the important information in a compressed format (= vector, "representation of the input")
- Decoder:
  - Can read the compressed representation from the encoder
  - Generates the output based on that information



# Sequence-to-sequence (seq2seq)

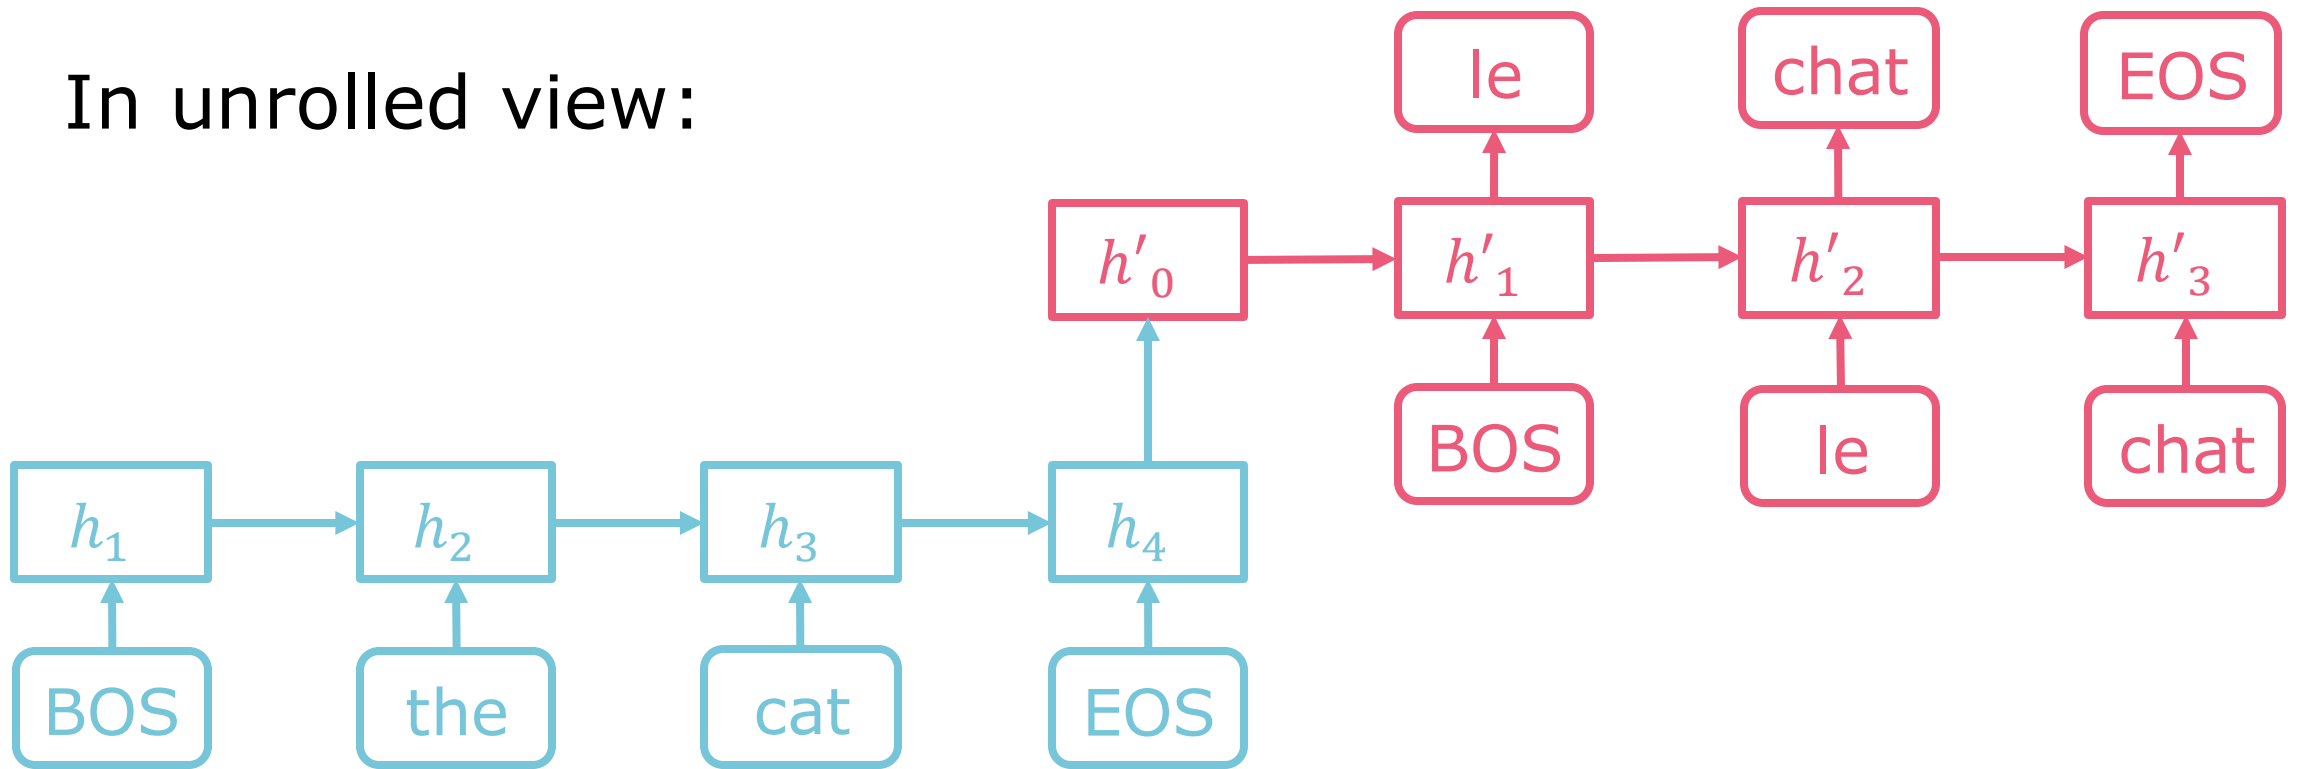
... or more like this:





# Sequence-to-sequence (seq2seq)

In unrolled view:



# Exercise: Seq2seq RNN

# Attention in RNNs

# Motivation

- Attention is the basis for Transformer (2017)
  - Start of a new era in NLP (and with a bit of delay also in other fields)
- The idea comes up in a machine translation paper written in 2014
- We will also see the first steps in pretraining (2018)
  - ... at least in NLP
  - computer vision was already routinely using features from models trained on ImageNet

# Human Attention

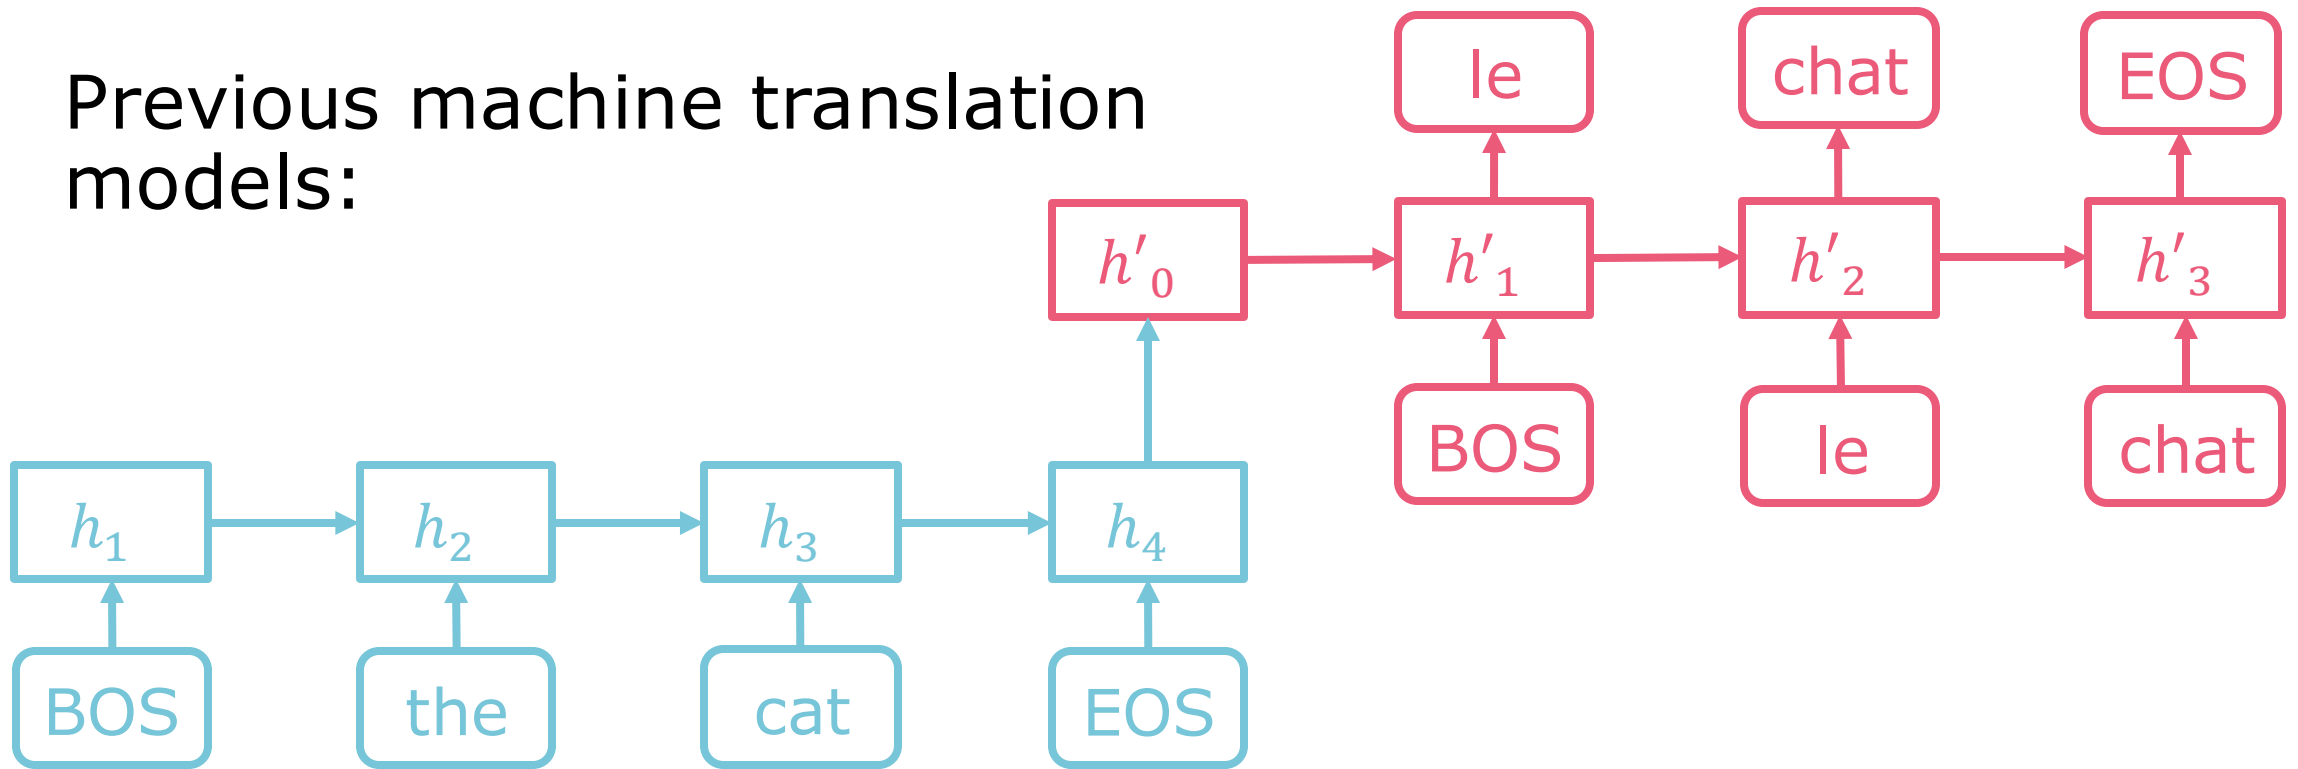
“Attention is the behavioral and cognitive process of selectively **concentrating on a discrete aspect of information**, whether considered subjective or objective, while ignoring other perceivable information.”

# Attention Motivation

- If we just use the final hidden state  $h_n$ , don't we throw away a lot of information computed in the earlier parts of the sequence?
  - Especially for long sequences: Information can be removed from the cell/hidden state at every time step  $t$ !
- Idea: Use a linear combination of all hidden states as the sequence's representation.

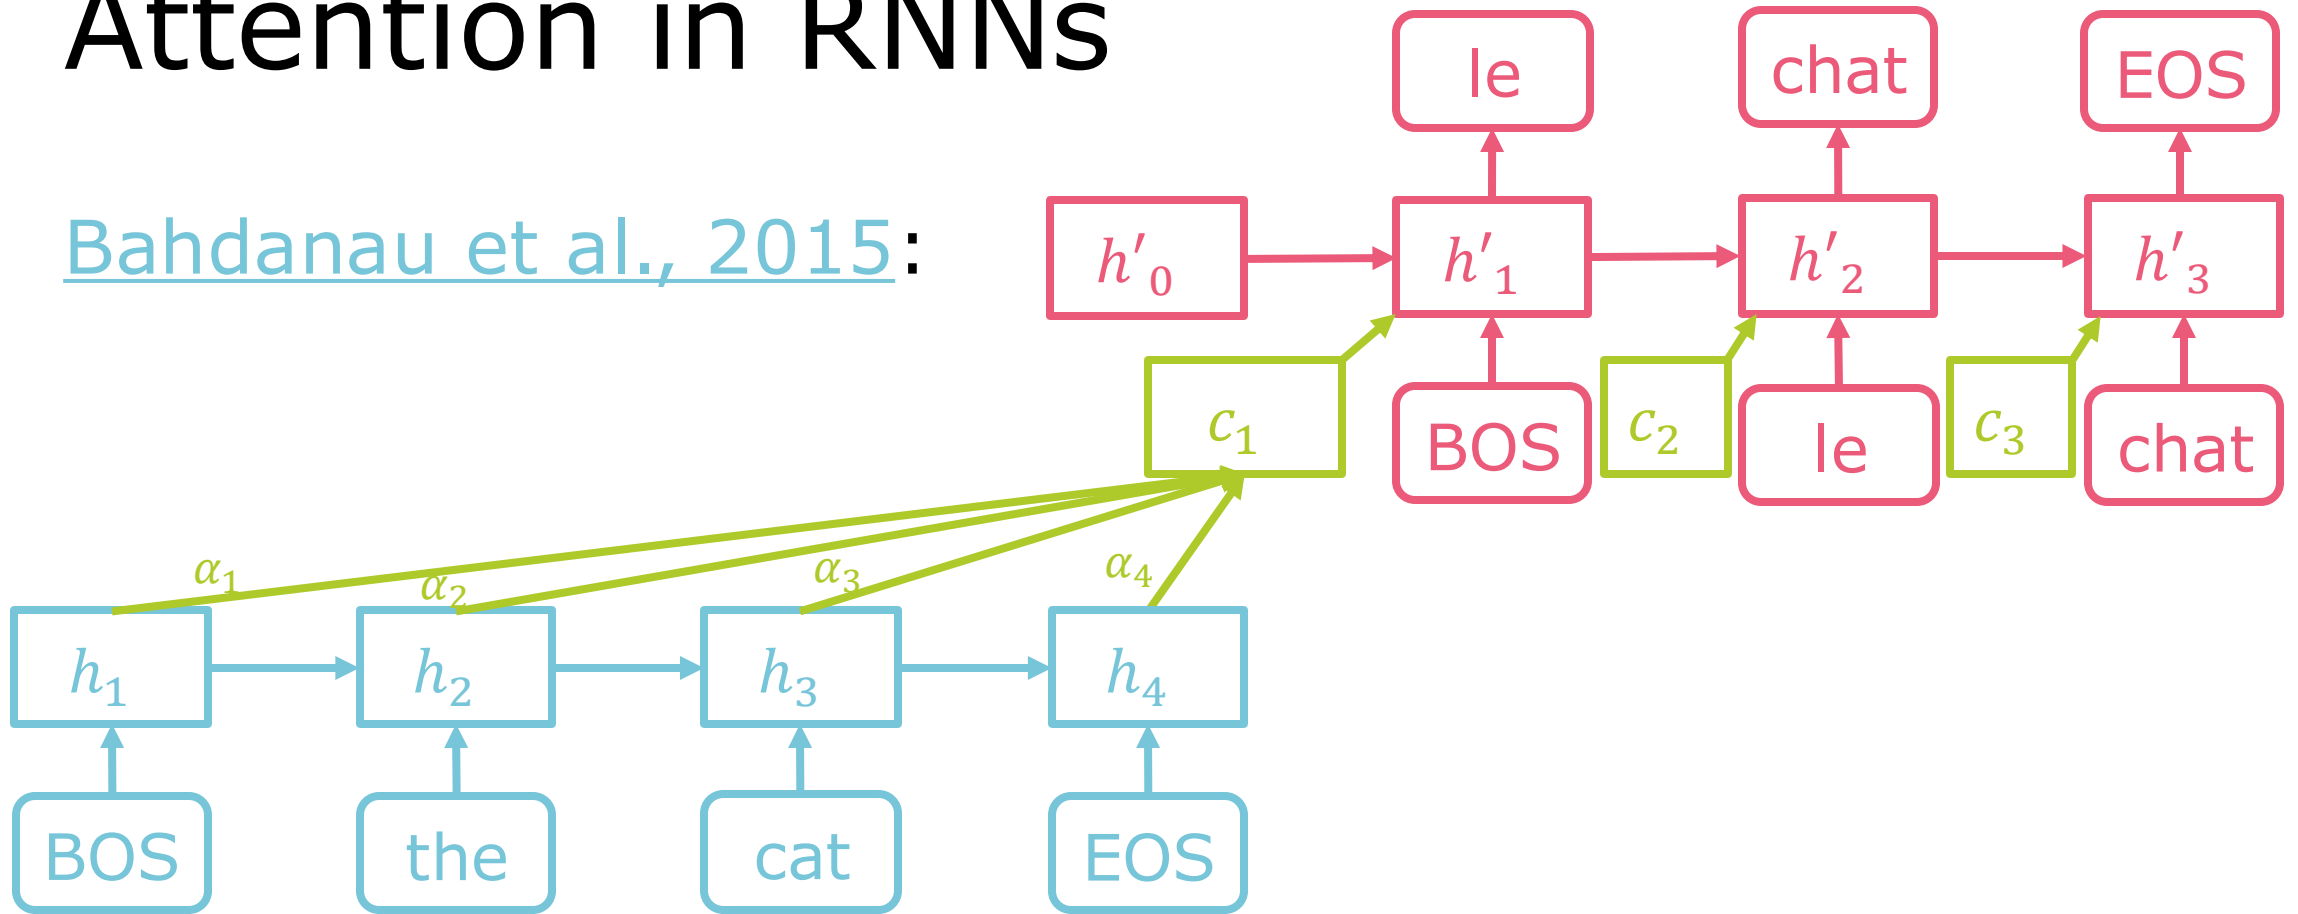
# Attention in RNNs

Previous machine translation models:



# Attention in RNNs

Bahdanau et al., 2015:





# Attention Definition

- Context vector  $c_i$ 
  - a weighted sum of encoder hidden states
  - computed for each decoding time step

$$c_i = \sum_{j=1}^n \alpha_{ij} h_j^{(\text{enc})}$$

i: time step in decoder  
j: time step in encoder

- Attention weights  $\alpha_{ij}$

$$\alpha_{ij} = \text{softmax}(e_{ij}) = \frac{\exp(e_{ij})}{\sum_{k=1}^n \exp(e_{ik})}$$

- Attention function  $f_{\text{att}}$

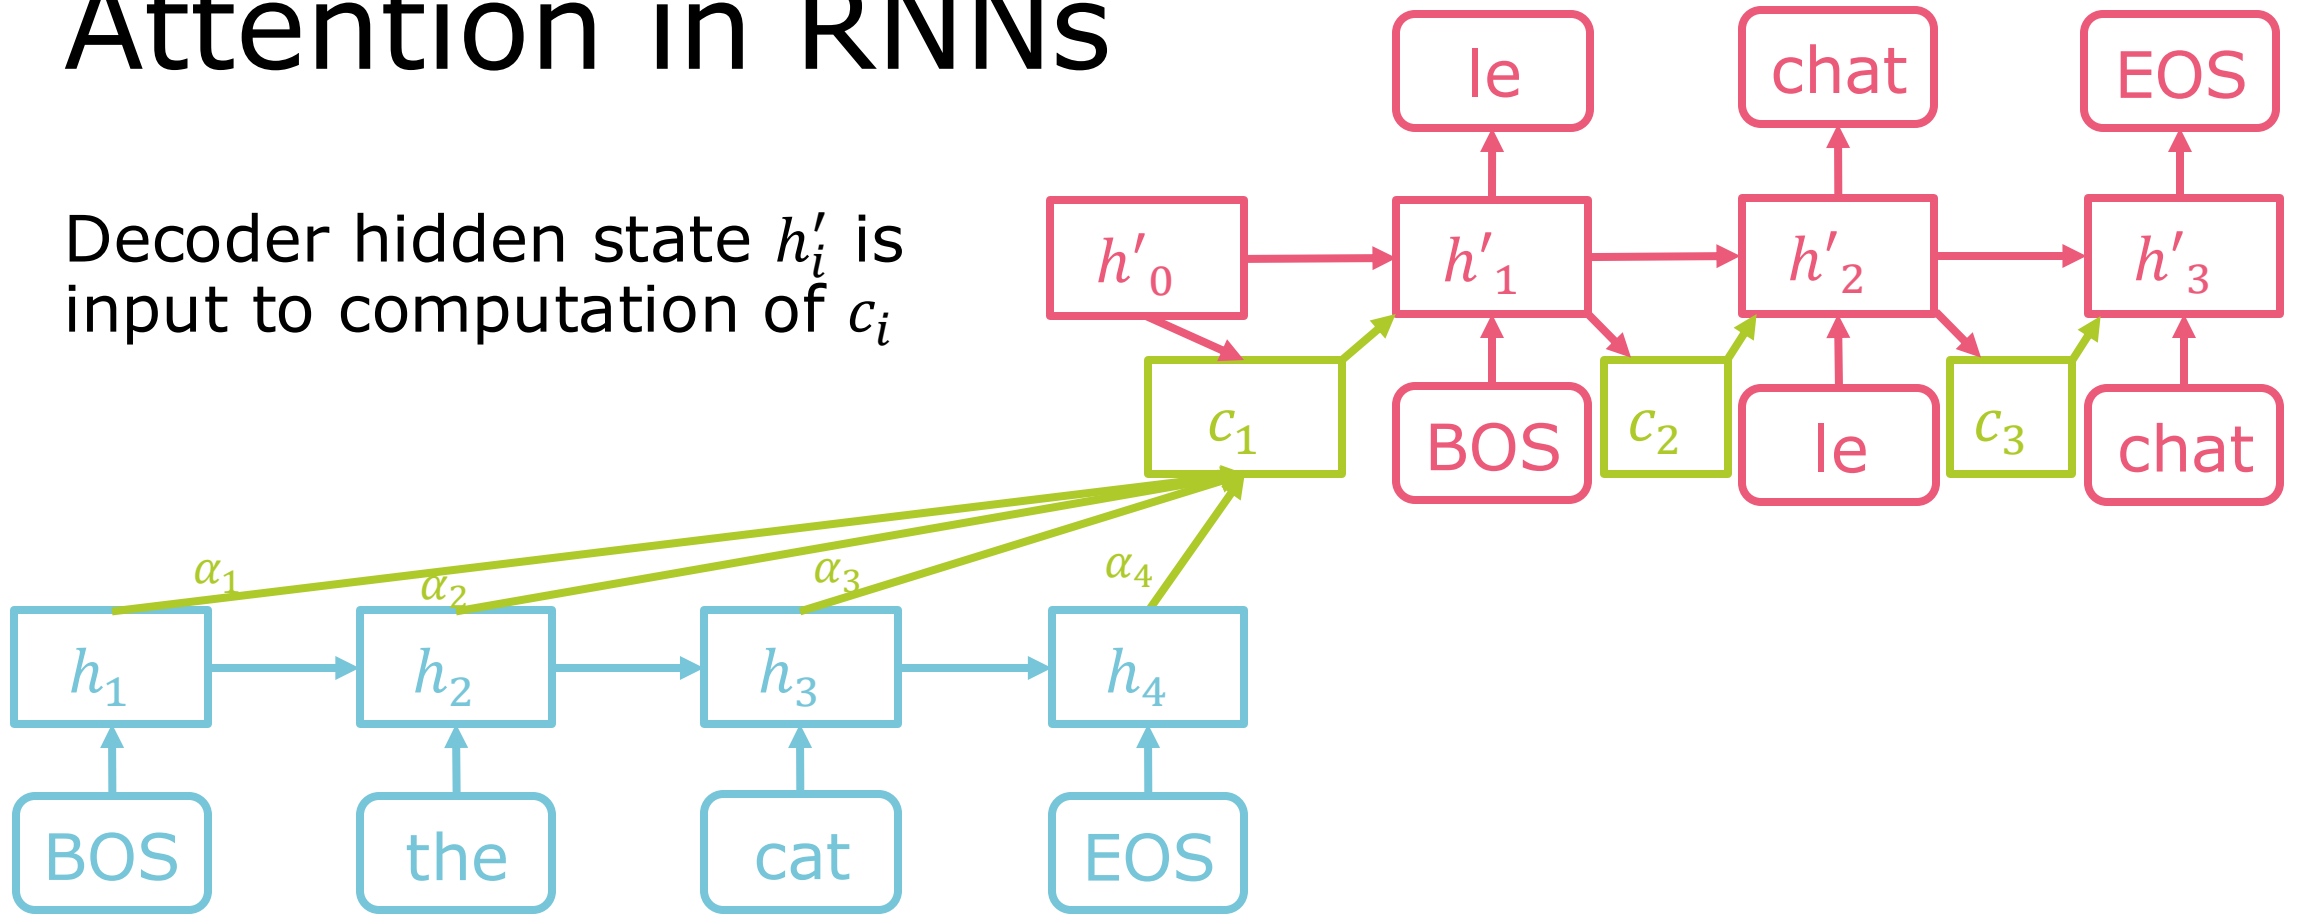
$$e_{ij} = f_{\text{att}}(h_{i-1}^{(\text{dec})}, h_j^{(\text{enc})})$$

previous decoder  
hidden state

current encoder  
hidden state

# Attention in RNNs

Decoder hidden state  $h'_i$  is input to computation of  $c_i$



# Attention Functions

- Additive attention

$$f_{\text{att}}\left(h_{i-1}^{(\text{dec})}, h_j^{(\text{enc})}\right) = v^{\top} \tanh\left(W h_{i-1}^{(\text{dec})} + U h_j^{(\text{enc})}\right)$$

- Multiplicative attention

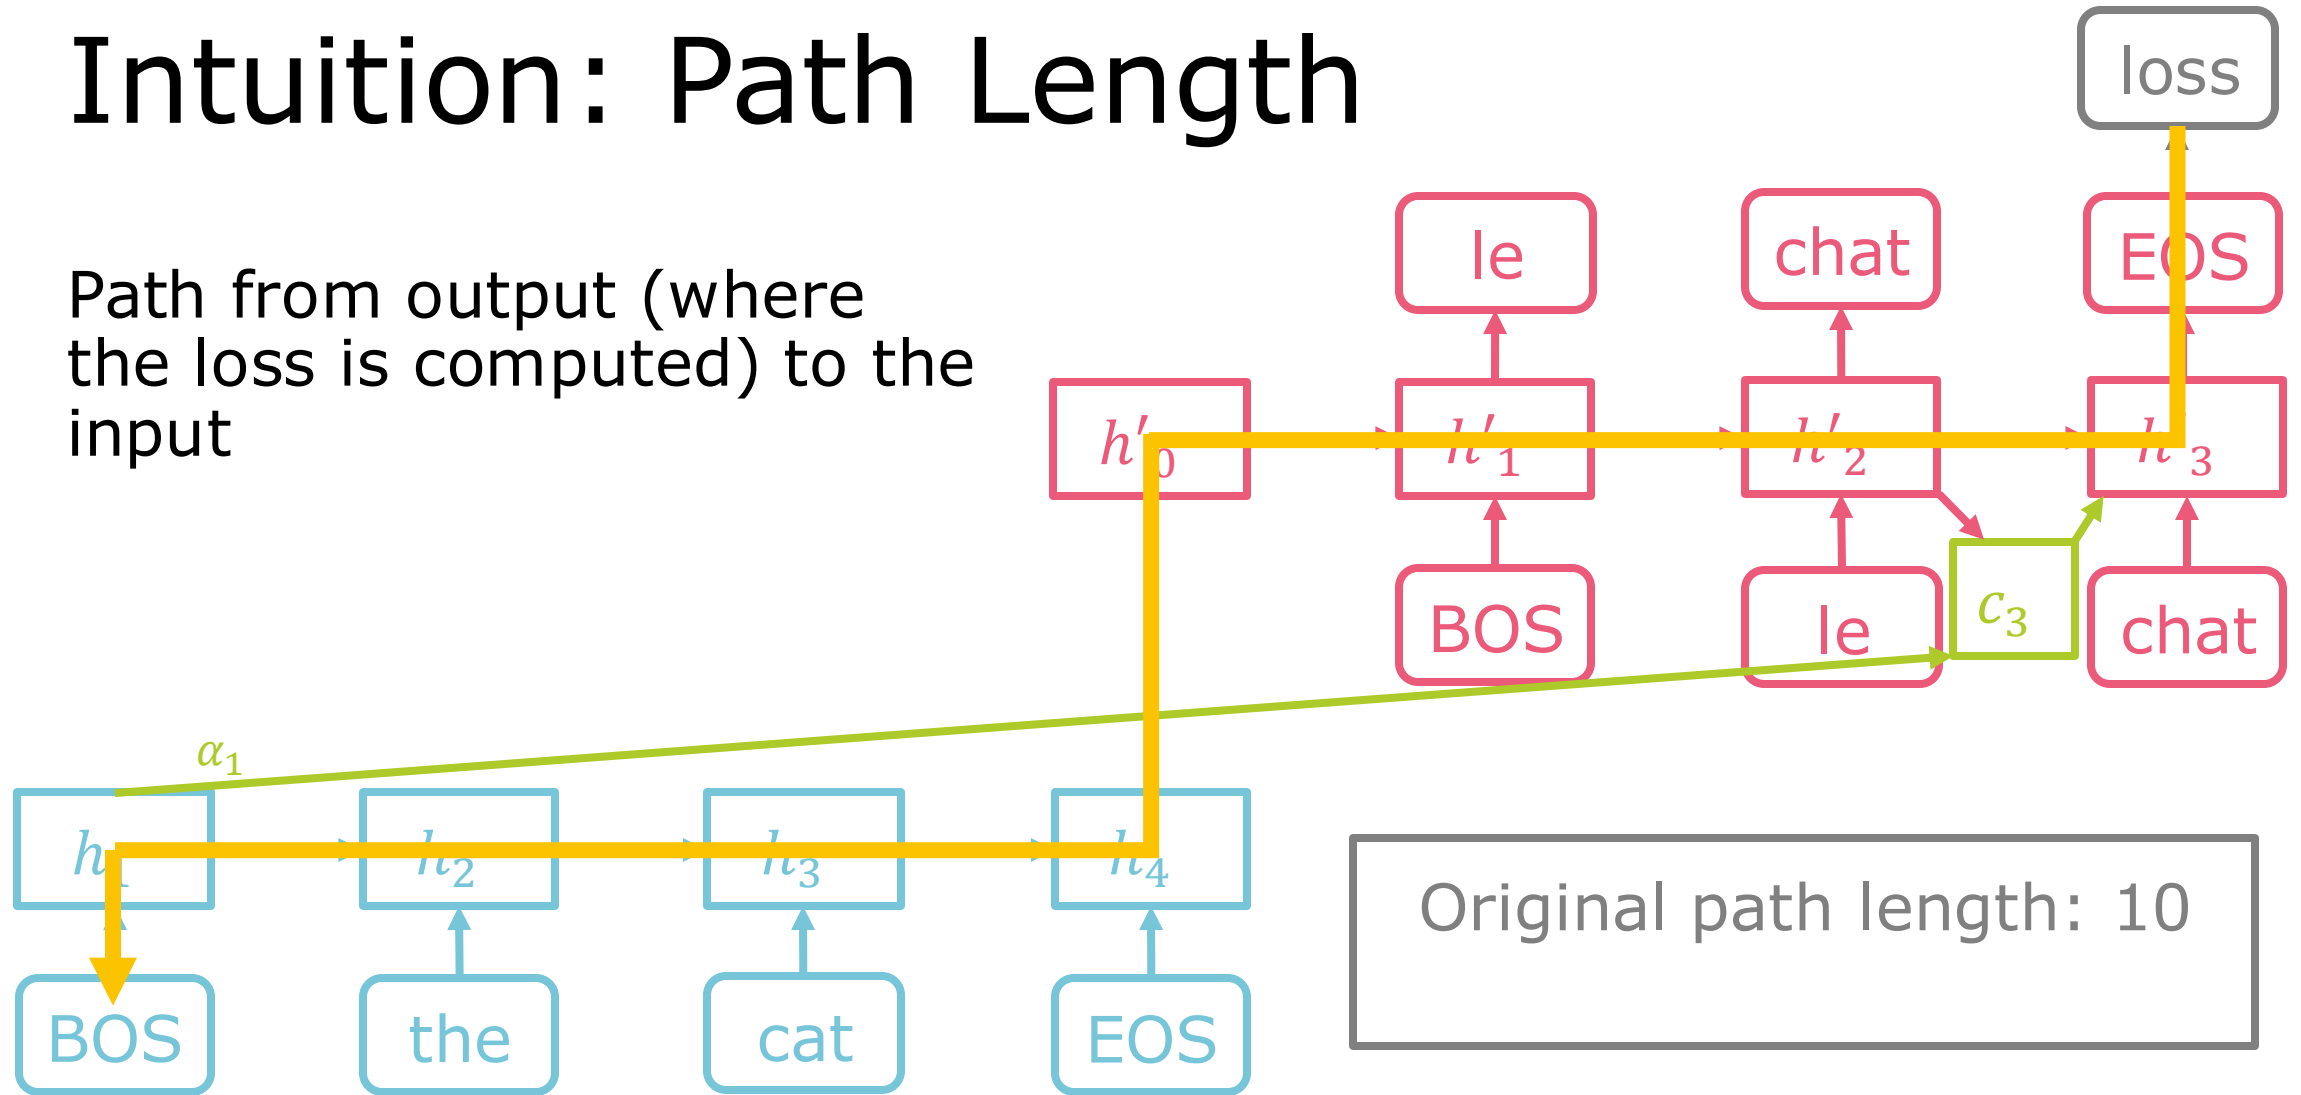
$$f_{\text{att}}\left(h_{i-1}^{(\text{dec})}, h_j^{(\text{enc})}\right) = h_{i-1}^{(\text{dec})} W h_j^{(\text{enc})}$$

learnable weights



# Intuition: Path Length

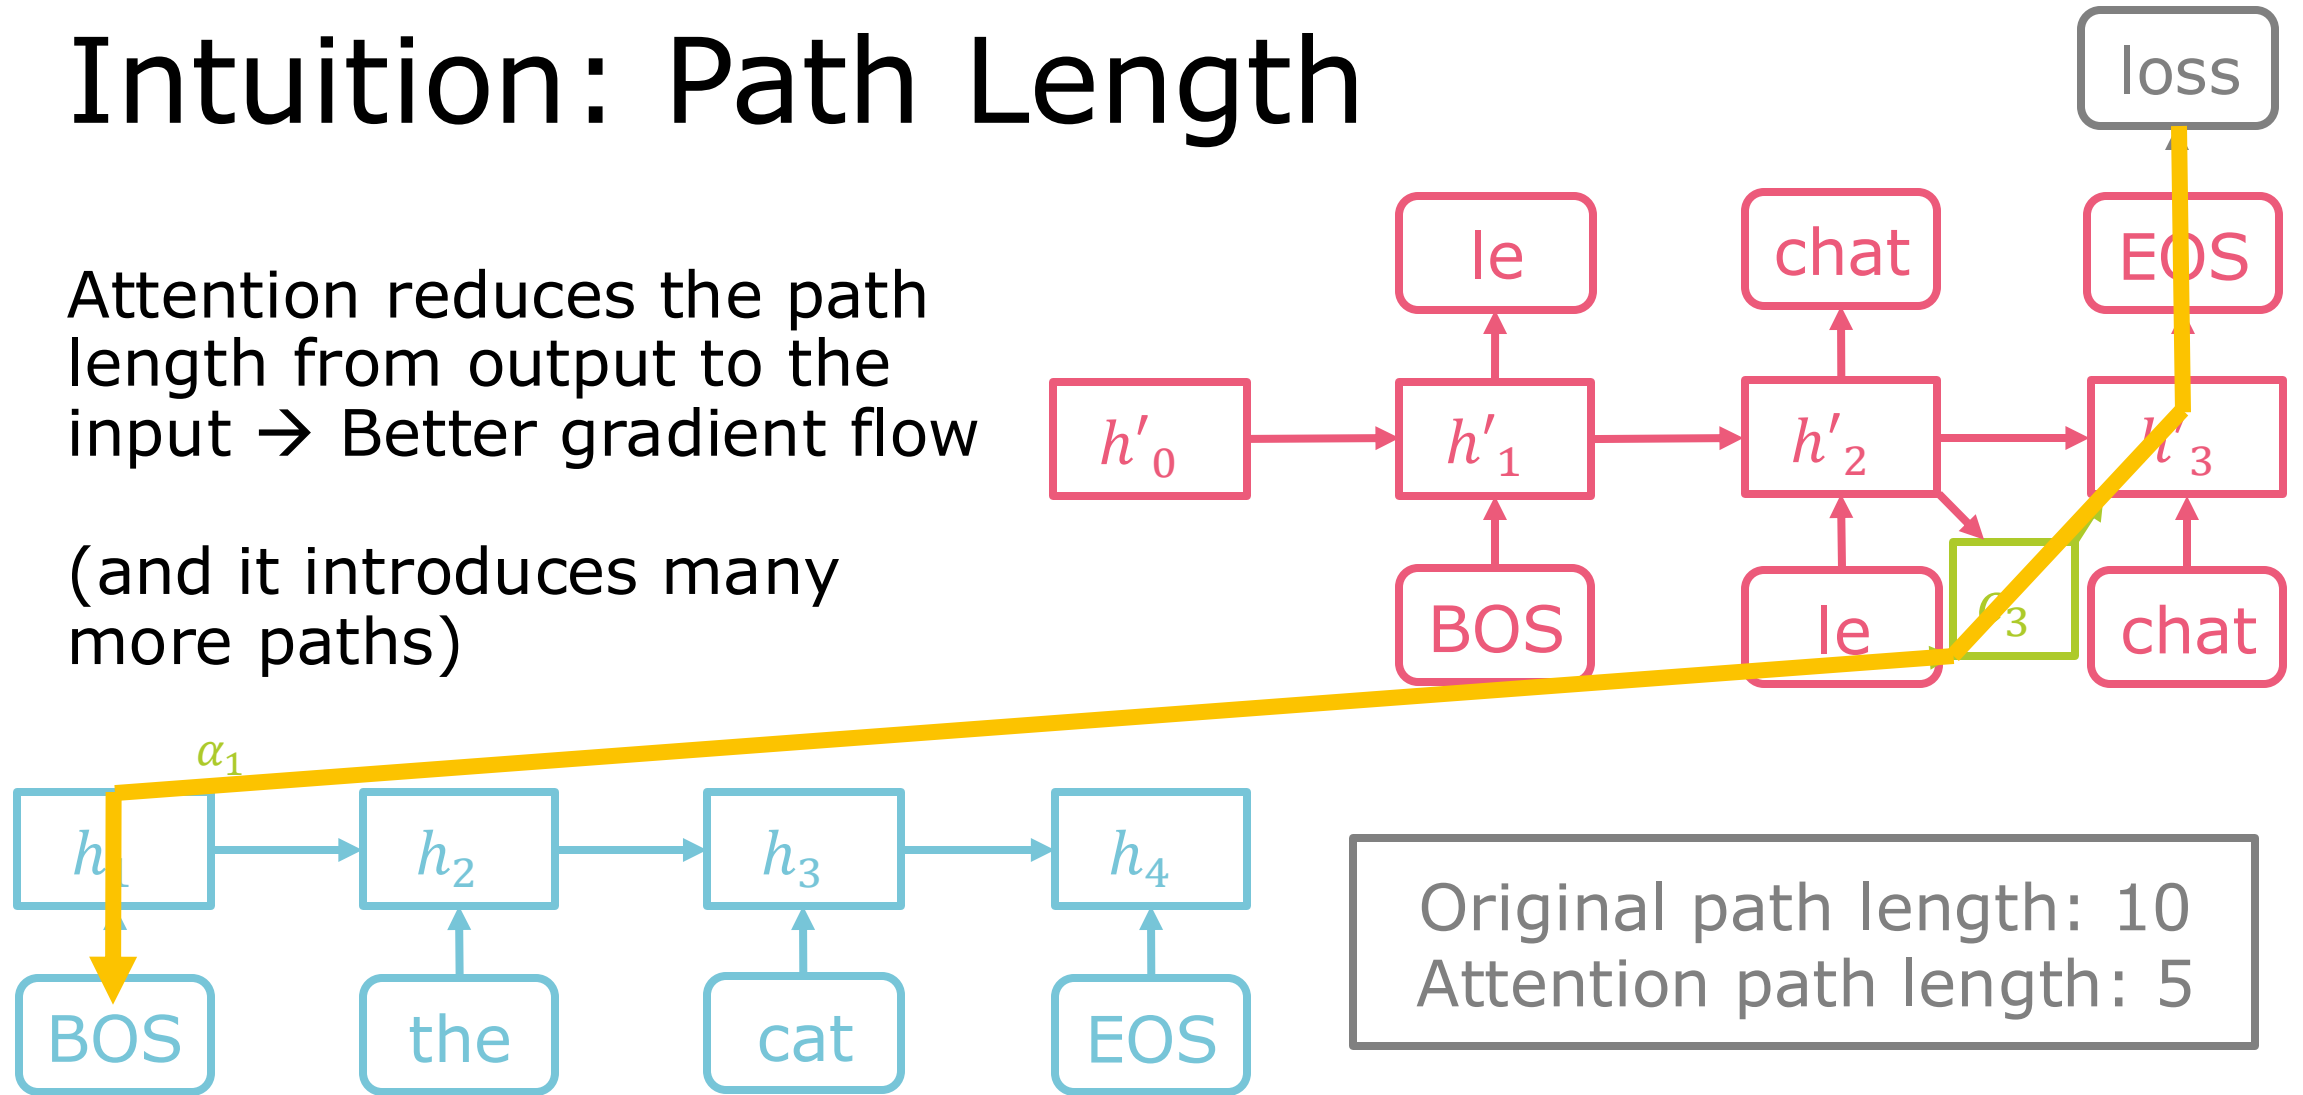
Path from output (where the loss is computed) to the input



# Intuition: Path Length

Attention reduces the path length from output to the input → Better gradient flow

(and it introduces many more paths)



# Exercise: Seq2seq RNN with Attention

# Contextual Embeddings: ELMo

# Embeddings from Language Models (ELMo)

**Deep contextualized word representations**

**Matthew E. Peters<sup>†</sup>, Mark Neumann<sup>†</sup>, Mohit Iyyer<sup>†</sup>, Matt Gardner<sup>†</sup>,**  
`{matthewp, markn, mohiti, mattg}@allenai.org`

**Christopher Clark\*, Kenton Lee\*, Luke Zettlemoyer<sup>†\*</sup>**  
`{csquared, kentonl, lsz}@cs.washington.edu`

<sup>†</sup>Allen Institute for Artificial Intelligence

\*Paul G. Allen School of Computer Science & Engineering, University of Washington



# Embeddings from Language Models (ELMo)

Mat



ardner<sup>†</sup>,

eg

,

u

\*Paul G. Al


of Washington

# Contextual(ized) Word Embeddings

- [Peters et al., 2018](#)
- Idea: Use the hidden states as embeddings of words in context
- word2vec & GloVe: static word embeddings, stays the same in each context
- ELMo: dynamic word embeddings, changes with the context

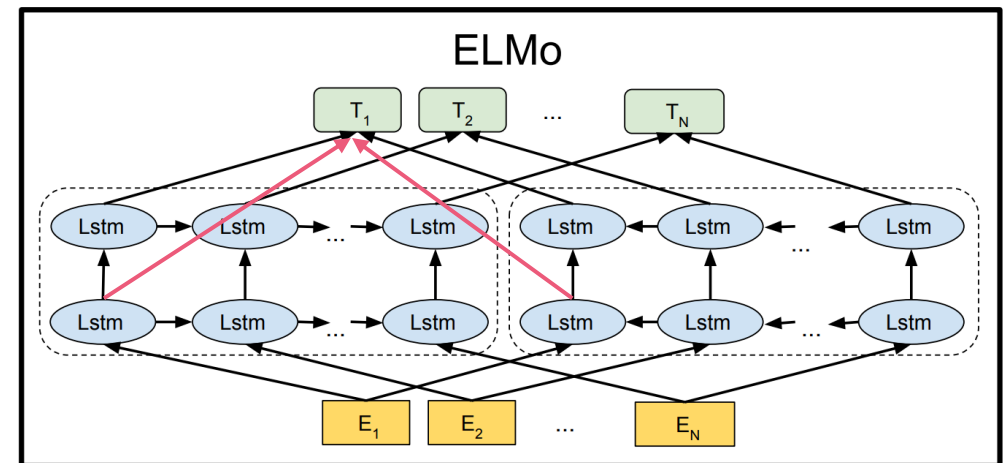
# Contextual Word Embeddings

- 2-layer bidirectional LSTM (BiLSTM)
- Word representation: Combination of LSTM's hidden states
  - Both directions
  - Both layers (imagine red arrows for each  $T_i$  in the image)



# Contextual Word Embeddings

- 2-layer BiLSTM
  - Layer 1 hidden states: Do well on syntax-related tasks (e.g. part-of-speech tagging)
  - Layer 2 hidden states: Do well on semantic (= meaning) tasks (word sense disambiguation)



# Contextual Word Embeddings

- 2-layer BiLSTM
  - Layer 1 hidden states: Do well on syntax-related tasks (e.g. part-of-speech tagging)
  - Layer 2 hidden states: Do well on semantic (= meaning) tasks (word sense disambiguation)

Using both layers performs better than just using layer 2


Task	Baseline	Last Only	All layers	
			$\lambda=1$	$\lambda=0.001$
SQuAD	80.8	84.7	85.0	<b>85.2</b>
SNLI	88.1	89.1	89.3	<b>89.5</b>
SRL	81.6	84.1	84.6	<b>84.8</b>

Table 2: Development set performance for SQuAD, SNLI and SRL comparing using all layers of the biLM (with different choices of regularization strength  $\lambda$ ) to just the top layer.

# Recipe for Contextual Embeddings

1. Train a BiLSTM on the 1B Word Benchmark for 10 epochs
2. Freeze the model's weights
3. For each example of your task:
  - a) Run the current input through the BiLSTM to get contextual word embeddings
  - b) Concatenate with the static word embedding and feed as input to the task network (in ELMo's case a different RNN)

# Recipe for Contextual Embeddings

1. Train a BiLSTM on the 1B Word Benchmark for 10 epochs
  2. Freeze the model's weights
  3. For each example of your task:
    - a) Run the current input through the BiLSTM to get contextual word embeddings
    - b) Concatenate with the static word embedding and feed as input to the task network (in ELMo's case a different RNN)
- 

# Task-specific Combinations

- We can combine the embeddings of each layer, depending on the task  
→ combine their strengths
- ELMo learns a scaling factor for each of
  - static word embedding
  - contextual embedding from L1
  - contextual embedding from L2
- Concatenate with input  $x_t$  ("input layer") or hidden state  $h_t$  ("output layer")




Figure 2: Visualization of softmax normalized biLM layer weights across tasks and ELMo locations. Normalized weights less than  $1/3$  are hatched with horizontal lines and those greater than  $2/3$  are speckled.



# ELMo Performance

- ELMo improves the state-of-the-art (SOTA) on various tasks

TASK	PREVIOUS SOTA		OUR BASELINE	ELMo + BASELINE	INCREASE (ABSOLUTE/ RELATIVE)
SQuAD	Liu et al. (2017)	84.4	81.1	85.8	4.7 / 24.9%
SNLI	Chen et al. (2017)	88.6	88.0	88.7 $\pm$ 0.17	0.7 / 5.8%
SRL	He et al. (2017)	81.7	81.4	84.6	3.2 / 17.2%
Coref	Lee et al. (2017)	67.2	67.2	70.4	3.2 / 9.8%
NER	Peters et al. (2017)	91.93 $\pm$ 0.19	90.15	92.22 $\pm$ 0.10	2.06 / 21%
SST-5	McCann et al. (2017)	53.7	51.4	54.7 $\pm$ 0.5	3.3 / 6.8%

# Nearest Neighbors for Contextual Word Embeddings

- Disambiguation of polysemous words (same word with multiple meanings)

	Source	Nearest Neighbors
GloVe	play	playing, game, games, played, players, plays, player, Play, football, multiplayer
biLM	Chico Ruiz made a spectacular <u>play</u> on Alusik 's grounder {...}	Kieffer , the only junior in the group , was commended for his ability to hit in the clutch , as well as his all-round excellent <u>play</u> .
	Olivia De Havilland signed to do a Broadway <u>play</u> for Garson {...}	{...} they were actors who had been handed fat roles in a successful <u>play</u> , and had talent enough to fill the roles competently , with nice understatement .

Table 4: Nearest neighbors to “play” using GloVe and the context embeddings from a biLM.